PREDICTING HOUSE PRICES USING MACHINE LEARNING

Loading the dataset:

The first step is reading the dataset from the csv file we downloaded.

```
import pandas as pd
import numpy as np


dataset = pd.read_csv("/content/USA_Housing.csv")


pd.options.display.float_format = '{:20.2f}'.format
dataset.head(n=5)
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.46	5.68	7.01	4.09	23086.80	1059033.56	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.64	6.00	6.73	3.09	40173.07	1505890.92	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	61287.07	5.87	8.51	5.13	36882.16	1058987.99	9127 Elizabeth Stravenue\nDanieltown, WI 06482...

let's get statistical information about the numeric columns in our dataset.

```
dataset.describe(include=[np.number], percentiles=[.5]).transpose().drop("count", axis=1)
```

	mean	std	min	50%	max
Avg. Area Income	68583.11	10657.99	17796.63	68804.29	107701.75
Avg. Area House Age	5.98	0.99	2.64	5.97	9.52
Avg. Area Number of Rooms	6.99	1.01	3.24	7.00	10.76
Avg. Area Number of Bedrooms	3.98	1.23	2.00	4.05	6.50
Area Population	36163.52	9925.65	172.61	36199.41	69621.71
Price	1232072.65	353117.63	15938.66	1232669.38	2469065.59

Then, we move to see statistical information about the non-numerical columns in our dataset:

Data Preprocessing:

Data cleaning:

Dealing with Missing Values: We should deal with the problem of missing values because some machine learning models don't accept data with missing values. Firstly, let's see the number of missing values in our dataset. We want to see the number and the percentage of missing values for each column that actually contains missing values.

```
num_missing = dataset.isna().sum()

num_missing = num_missing[num_missing > 0]

percent_missing = num_missing * 100 / dataset.shape[0]

pd.concat([num_missing, percent_missing], axis=1,
          keys=['Missing Values', 'Percentage']).\
          sort_values(by="Missing Values", ascending=False)
```

Missing Values	Percentage

```
dataset["Price"].value_counts()

    1059033.56    1
    1521141.34    1
```

```
1148372.40    1
2065710.16    1
1749820.01    1
             ..
1444701.33    1
 788427.84    1
 875904.53    1
 984421.23    1
1298950.48    1
Name: Price, Length: 5000, dtype: int64
```

```
dataset['Price'].fillna('No feature', inplace=True)
```

let's check if there is any remaining missing value in our dataset:

```
dataset.isna().values.sum()
```

```
0
```

Converting categorical features into numerical representations:

We will encode our categorical features using one-hot encoding technique which transforms the categorical variable into a number of binary variables based on the number of unique categories in the categorical variable.

```
dataset[['Address']].head()
```

	Address
0	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	9127 Elizabeth Stravenue\nDanieltown, WI 06482...
3	USS Barnett\nFPO AP 44820
4	USNS Raymond\nFPO AE 09386

```
dataset = pd.get_dummies(dataset)
```

```
address_oneHot = [c for c in dataset.columns if c.startswith("Address")]
dataset[address_oneHot].head()
```

	Address_000 Adkins Crescent\nSouth Teresa, AS 49642-1348	Address_000 Todd Pines\nAshleyberg, KY 90207-1179	Address_001 Steve Plaza\nJessicastad, UT 25190	Address_0010 Gregory Loaf\nSouth Ericfort, VA 34651-0718	Address_00 Raym Knolls\n Jason, 75
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	

5 rows × 5000 columns